QC Spec Doc - V2

Changelog

Last Updated: Apr 16, 2026

Date	Type	Description
Apr 16, 2026		*Search **04/15*** - *Universe date correctness updates and exceptions*
Apr 13, 2026		*Search **04/11*** - *Minor changes regarding the Keystone universe* - *New personas and business functions added*
Apr 9, 2026		*Search **04/09*** - *Added "All-Failing Rubrics" dimension to flag invalid rubrics that all the completed runs fail on* - *Added [Non-Fail - Outcome-Light] to flag when outcome rubrics are more than 50% rubrics, but less than 80%*

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

Date	Type	Description
Apr 2, 2026		*Search **04/02*** - *CBs now use the new "Scenario Generation" tool to make changes to the base universe to create "Enhanced Universes", which can then be edited using SQL queries like before* Note: Scenario Gen is offline since Apr 07.
Mar 29, 2026		*Search **03/29*** - *New: "Fail – Too Many Erroneous Runs" status added.* - ⚠️ *This is different from runs that fail because of rubric issues.*
Mar 26, 2026		*Search **03/26*** - *New Rubric Dimension: "Overly Broad Criteria" — catches criteria that are too vague or general.* - *Added a small note comparing Incorrect Criteria with Overly Broad Criteria*

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

Date	Type	Description

Mar 24, 2026		*Search* ***03/24***
		• *The universe now has a fixed date for "today" – April 26, 2026*
		  ○ *Added this as a workflow note*
		  ○ *Added a new rubric dimension "Alignment with Today's Date" to make sure that the prompt and universe make sense considering "today's" date*

Note: Ensure your submissions meet all Quality Check (QC) rubrics. This document provides a complete list of the criteria used by QC to audit and determine if a task passes or fails.

Scoring Guidelines

Dimension	Sub-Dimension	Notes for Auditors	1 - 2	3 - 4	5
Prompt	**Unique Ground Truth**	The prompts are open-ended in this project, so there cannot be one possible response to the user (For example, wording can differ), BUT it should not be the case that the prompt leads to *distinct*	**[Fail - Multiple Valid Answers]** - The prompt can reasonably produce two or more different correct answers, and the prompt does not clearly indicate which answer is intended.	**[Non-Fail - Multiple Valid Answers]** The prompt is ambiguous, but there is either: a leading interpretation that the majority of people would pick if polled,	- The prompt has a unique ground truth answer. All the experts on the subject would come to the same conclusion

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

		outcomes. It's okay if there's multiple possible paths the agent could take to arrive at the answer, but the key points in the answer, as requested in the prompt, should be the same.		or there are two different valid interpretations where making a small, logically defensible assumption would clarify which was intended by the user.	
Prompt	**Feasibility**		**[Fail - Prompt Impractical Request]** - Prompt contains an impractical request that can't be answered by an LLM in a single response **[Fail - Prompt Impossible Request]** - Prompt has at least one request that can't be fulfilled at all **[Fail - Prompt Conflicting Instructions]** - Prompt gives instructions that conflict/contradict with itself that can't be fulfilled simultaneously **[Fail - Prompt Feasibility with Tools]** - Prompt contains an impractical or impossible request that can't be answered by the tools available or enabled for the task	**[Non-Fail - Prompt Impractical Secondary Request]** - There are multiple requests in the prompt and it's verging on being impractical for a model to answer it in a single response, but the core request of the prompt can be fulfilled with minor concessions on the secondary requests	- The prompt is completely actionable by an LLM agent. - The prompt contains no conflicting instructions/statements
Prompt	**Explicit Tool Mention** **UPDATED 03/13**		**[Fail - Explicit Tool Mention]** - The prompt explicitly mentions an MCP tool name (e.g., linear_create_issue, crm_search_contacts, quickbooks_create_bill) **03/13** ~~The prompt explicitly mentions a specific tool name~~ ~~The prompt unnaturally refers to an MCP server in an inorganic, technical~~	**03/13** **[Non-Fail - Tool/MCP Server Mentioned]** - The prompt references an MCP server by name (e.g., "use the Linear MCP server", "via the Gmail MCP server") - The prompt refers to any system as a "tool" (e.g., "use the QuickBooks tool", "use the CRM tool")	- The prompt does not explicitly state the use of tools

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

			~~manner~~ ~~"Use the Airbnb tool to…" (technical, unnatural)~~ ~~"I want to use Airbnb to find…" (general, natural)~~		
Prompt	**Prompt Clarity and Specificity**		**[Fail - Major Clarity / Specificity Issues]** - It's not clear what is being asked, the prompt is extremely difficult to follow, or is overly vague; the user intent cannot be reasonably understood from the prompt. - Major details are missing that are needed to answer the prompt and cannot be reasonably assumed	**[Non-Fail Minor Clarity / Specificity Issues]** - It's mostly clear what is being asked but the request could reasonably be interpreted multiple ways	- There is little to no room for misinterpretation of the specific request - Prompt has a specific request that doesn't require more than one minor assumption to answer it

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

Prompt	**Contrived / Unnatural Prompts**	This should check for artificial difficulty vs. natural complexity. Contrived examples include: demanding exact timestamps ("Find the email from January 15th at 3:47 PM"), requiring arbitrary format constraints ("Respond in exactly 3 sentences using passive voice"), or asking the agent to do something no real employee would do ("Intentionally send an incorrect email to test error handling"). Natural difficulty comes from entity confusion, scattered information, conflicting data across services, and hidden root causes. Prompts written as step-by-step command lists ("First search emails, then check CRM…") should also be considered contrived.	**[Fail - Contrived / Unnatural Prompt]** - Prompt is contrived or otherwise unnatural note: prompts such as riddles that are intentionally constraints/details laden should not be considered contrived.	**[Non-Fail - Somewhat Contrived / Unnatural Prompt]** - Prompt is somewhat contrived or is somewhat unnatural note: prompts such as riddles that are intentionally constraints/details laden should not be considered contrived.	- The prompt is neither contrived or unnatural
Prompt	**Truthfulness**		**[Fail - Major Factual Errors]** - The prompt contains 1 or more major factual errors **[Fail - Minor Factual Errors]** - The prompt contains 2 or more minor factual errors (e.g., statements irrelevant to the fulfillment of the prompt's request)	**[Non-Fail - Minor Factual Errors]** - The prompt contains 1 minor factual errors.	- The prompt has no factual errors - The prompt has no misleading statements

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

Prompt	**Tool use and Cross-service requirement**		**[Fail - Tool Dependence]** - The prompt can be answered without having to make any tool calls (either through basic reasoning or general knowledge) **[Fail - Cross-service Requirement]** - Prompt can be fully resolved within a single service (e.g., just email or just CRM), or only trivially touches a second service.	N/A	- Prompt requires investigation across 2+ services to answer correctly (e.g., facts needed are scattered and must be reconciled).
Prompt	**Investigation UPDATED 03/08**		**03/08** ~~[Fail - Action Missing]~~ ~~- Prompt only asks for a summary/report/research with no required action.~~ **[Fail - Pre-solved]** - Prompt tells agent the root cause and no investigation is needed.	N/A	Prompt clearly requires investigation ~~and at least one concrete action (email, ticket creation/update, booking, calendar scheduling, CRM update, etc.)~~ **03/08**, and it logically follows from investigation.

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

Prompt	**Coherence**	An example of failing bolt-on requests: "Check weather in Miami, update my calendar for next week, email Marcus about Q3, and look up Seattle flights." Why it fails: Four unrelated requests stapled together. Each sub-request should causally flow from the same situation. How to tell if your prompt is "bolted": Remove any one sentence. If the rest still makes perfect sense without it, that sentence is a bolt-on	**[Fail - Bolt-on requests]** - Contains unrelated bolt-on requests; removing a sentence doesn't change the rest.	N/A	One cohesive situation; stacked asks all tie back to the same purpose.
Prompt	**Persona**	See the list of personas in the instructions doc or in the summary doc **04/11 UPDATE**: Persona List - v2.2	**[Fail - Persona Mismatch]** The prompt written by the CB does not read like it could have been written by the persona selected. **03/24 NOTE**: Do not penalize with a failure if the prompt aligns with multiple Personas, so long as the assigned Persona is also applicable for the prompt.	**[Non-Fail - Persona Mismatch]** The prompt written by the CB is plausible for the persona selected, but fits a different persona better	The prompt written by the CB aligns with the persona selected
Prompt	**Business Function** **ADDED 03/20**	Each task is assigned a business function (Operations, Customer Engagement, Engineering, Finance, Executive). The prompt should reflect a realistic scenario that falls within that business function.	N/A	**[Non-Fail - Business Function Mismatch]** The prompt's primary scenario is ambiguous — it could reasonably belong to either the assigned function or a different one.	The prompt clearly reflects a realistic scenario within the assigned business function.

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

		List of 5 business functions in this universe here: https://gist.github.com/sunjiehou-scale/f242d21d24f5544bdc84dd3199a1094e This is only applicable to v2.1 tasks. Each task version would either be v2 or v2.1 mentioned at the top of the task 04/11 **UPDATE**: Business Functions - v2.2			
Prompt and Universe	**Alignment with Today's Date** **ADDED 03/24**	**NOTE:** The universe now has a **fixed date for "today" – Apr 26, 2026 (v2.1).** This date it fed to the agent in the system prompt (we can't see the system prompt) The prompt and universe should make sense in the context of this date. 04/11 **NOTE**: The fixed date for the **v2.2** universe is Apr 28, 2026. 04/15 **NOTE**: Universe not aligned with the correct date should not fail if it did not cause an agent failure.	**[Fail - Prompt Misaligned with Today]** - The prompt's request does not make sense or can't be fulfilled, given "today's" date **E.g.,** "What happened in May 2026 that the sales are down 30% month over month?" is nonsensical because it's asking about the future **[Fail - Universe Misaligned with Today]** - CB modified the universe in a way such that it's incoherent with "today's" date. For example, adding a new Slack message sent on May 1, 2026. - The default universe data is not aligned with "today's" date..AND it caused an agent failure. 04/15 **Note:** Events that are meant to be in the future are still okay to be present. For example, a calendar event scheduled to	N/A	Prompt does not contradict "today's" date The default universe data is not aligned with "today's" date..BUT it did not cause an agent failure. 04/15

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

			be in May is okay because it's meant to be future facing		
Universe	**Universe Feasibility (Data Exists)**	Auditors should check the Universe Explorer or Sandbox to verify that key data needed for the prompt exists.	**[Fail - Missing Universe Data]** - Task depends on key facts/entities/events that do not exist in the universe (or cannot be retrieved with available tools), making the task unsolvable.	N/A	All core facts required to solve the task exist in the universe and are retrievable via tools.
Universe	**Cross-service Coherence**	See the audit workflow to understand how CBs modify the universe and how you can view the changes they made **(04/02)**	**[Fail - Incoherent Universe Changes]** - Universe edits create contradictions that break solvability or realism (e.g., mismatched names across email vs CRM, impossible dates, entities referenced in one service but not elsewhere when required). **Note**: When evaluating a contradiction, if there is sufficient supporting evidence to support one piece of information against a contradicting low-supported piece of information, do not count that as a contradiction.		Universe edits (if any) are internally consistent and coherent across services, including dates, names, and storyline continuity.

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

Oracle Events (OEs)	**OE Completeness**	Additional Notes: OEs are the key steps a correct agent would take. They serve two purposes: (1) prove solvability, and (2) drive rubric writing across all three categories — Outcome, TS, and QC **(03/20)** Each OE should note: what action happens, what info is discovered, what tool is used, and what parameters matter. OEs describe tool-use steps — not what the final response should say.	N/A	**[Non-Fail - Incomplete Oracle Events]** - OEs are missing critical steps needed to solve the task (e.g., no step describing how key info is retrieved, or how required action is executed).	OEs describe the full critical path: key discovery steps + dependency chain(s) + required write action(s).
Oracle Events (OEs)	**OE Accuracy**	Auditors should verify OEs against the universe data using the Explorer or Sandbox. Each OE should reference the correct tool, correct service, correct parameters, and correct expected data. An OE that describes "look up Fatimah's relocations in Linear" when relocation records are actually in the CRM is inaccurate — the data exists but the OE points to the wrong place. OE Accuracy is distinct from OE Completeness: completeness checks whether steps are missing; accuracy checks whether the steps that exist are correct.	N/A	**[Non-Fail - Minor OE Inaccuracies]** OEs are substantively correct but contain minor imprecisions (e.g., slightly wrong parameter name that doesn't change the outcome, approximate rather than exact expected values for freetext fields). **[Non-Fail - Inaccurate Oracle Events]** One or more OEs reference the wrong tool, wrong service, wrong parameters, or wrong expected data. The described steps would not produce the correct results if followed literally (e.g., searching the wrong service for data that lives elsewhere, expecting data that doesn't match the universe, specifying incorrect entity names or identifiers).	All OEs are factually accurate. Tools, services, parameters, and expected data match the universe. Following the OEs literally would produce a correct trajectory.

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

Rubrics	**Overall Rubric Quality**		Use the number of criteria that the CB wrote as the denominator while calculating % values. See the additional notes section for the numerator. Do NOT double count criteria while tallying even if it has multiple issues. **[Fail - 10%+ Major Rubric Errors]** More than 10% (>10%) of the criteria contain major issues **[Fail - 15%+ Moderate Rubric Errors]** More than 15% (>15%) of the criteria contain moderate or major issues **[Fail - 20%+ Minor Rubric Errors]** More than 20% (>20%) of the criteria contain minor or moderate or major issues	**[Non-Fail - Up to 10% Major Errors]** Up to 10% (<=10%) of the criteria contain major issues **[Non-Fail - Up to 15% Moderate Errors]** Up to 15% (<=15%) of criteria contain moderate or major issues (with major issues contributing lower than 5%) **[Non-Fail - 5-20% Minor Errors]** Between 5 and 20% (>=5% and <=20%) of criteria contain minor or moderate or major issues (with major issues contributing lower than 5% and moderate issues contributing lower than 15%)	- Less than 5% (<5%) of the rubrics have minor issues - No major or moderate issues
Rubrics	**All-Failing Rubrics** **ADDED 04/09**	After rubric verifiers run, CBs are required to review any rubric that failed all completed runs (namely “All-failing” rubrics or AF rubrics) and add a 1–2 line justification confirming the failure is a genuine model failure and not a rubric quality issue. "All completed runs" means only runs that finished without an error — if 4 of 6 completed, the rubric must have failed all 4. If none of the rubrics failed all completed runs, this step is empty and no justification is needed. In some tasks, this will be in a text box	**[Fail - 2+ Invalid All-Fail Rubrics]** Two or more all-failing rubrics have a quality issue that causes the failure — e.g., fictional tool name, overly specific wording, beyond prompt scope, unfairly penalizes a valid agent approach, or not grounded in GT. The CB either did not notice the rubric problem or provided a justification that inadvertently reveals it. This indicates a false negative — the task may appear harder than it is due to a bad rubric.	**[Non-Fail - Poor All-Fail Rubric Justification]** The all-fail rubric is good quality (i.e., the model failures for that rubric in all the runs are genuine fails), but the explanation provided by the CB is either unfactual, irrelevant or doesn’t support the validity of the all-failing rubric. **[Non-Fail - One Invalid All-Fail Rubric]** Exactly one all-failing rubric has a quality issue that causes the failure — e.g., fictional tool name, overly specific wording, beyond prompt scope, unfairly penalizes a valid agent approach, or not	All rubrics that failed all completed runs are valid — self-contained, grounded in GT, within prompt scope, and using real tools. The CB's justifications confirm the failures reflect genuine model reasoning or execution gaps, not rubric artifacts. If no rubrics failed all completed runs, this dimension is automatically a 5.

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

		at the end of the rubric run step and in most new tasks (this is the format we're backfilling tasks to), this will be in the rubric rating step at the end of the verifier runs. If no AF rubrics exist, this dimension doesn't need to be audited.		grounded in GT. The CB either did not notice the rubric problem or provided a justification that inadvertently reveals it. This indicates a false negative — the task may appear harder than it is due to a bad rubric. (**Note:** There may be more all-failing rubrics, but only one of them has to have a quality issue for the non-fail to apply. If more all-failing rubrics have quality issues that cause failure, then the task fails)	
Rubrics	**Rubric Category Balance**	Outcome rubrics check final correctness (was the result right?). Process rubrics = Tool Selection + Query Construction (did the agent use the right tools with the right parameters?). Target ratio is approximately 80% Outcome, 20% Process.	**[Fail - No Outcome Rubrics]** - The rubric set contains zero Outcome rubrics. **[Fail - Process-Heavy]** - More than 50% of rubrics are **Process** (Tool Selection + Query Construction), with **Outcome** rubrics underrepresented. ~~[Fail - Missing Grounding Check] Investigation-heavy task has no grounding check rubric verifying the agent retrieved data before reporting it.~~ **(03/20)**	**[Non-Fail - Outcome-Light]** - **Outcome** criteria account for at least 50%, but less than 80% of the criteria present. Process Rubrics (TS + QC) account for 50% or less. **(04/09)**	80% or more of rubrics are **Outcome**. **Process** rubrics (TS + QC) are present but limited to critical oracle-event steps. **(04/09)**
Trajectory	**Tool Call Count** **UPDATED 03/20**		**[Fail - Insufficient Tool Calls]** - The average tool call count across all trajectories is less than 15		The average tool call count across all trajectories is greater than or equal to **03/29** 15.

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

Trajectory	**Agent Failure Rate**	At most 40% of the runs should pass the rubrics. If pass@1 > 40%, the task would fail. Given there are supposed to be 6 runs, this translates to at least 4 of the agent runs needing to fail the prompt. **TEMPORARILY REMOVED 03/08** ~~**A failure is defined as the model failing at least one valid outcome rubric. Process rubric failures are not model failures.**~~ Note: You may come across tasks where the rubric verifiers did not run successfully (error or runs indefinitely). In such cases, if you are able to validate that the failed runs that did finish are true failures, and this happens to be at least 60%, that's sufficient for satisfying this dimension. Otherwise, we might need to manually validate the trajectories to see if they failed.	**[Fail - Too Easy]** - More than 2 of 6 runs pass all valid ~~outcome~~ **03/08** rubrics (i.e., pass@1 > 40%)		0–2 of 6 runs pass all ~~outcome~~ **03/08** rubrics (pass@1 ≤ 40%), proving the task is sufficiently difficult
Trajectory	**Error Rate** **ADDED 03/29**	Sometimes, runs may error out. CBs are allowed to continue as long as there are at least 4 successful runs (i.e., max of 2 erroneous runs). **NOTE**: This is not the same as a "run with rubric failures".	**[Fail - Too Many Erroneous Runs]** - At least 3 runs were not successfully completed (resulted in an Error).	N/A	- Fewer than 3 runs were not successfully completed (resulted in an Error).

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

Appendix

Rubric Quality Definitions

Major Issues	**Missing Criteria — Outcome**	**Definition:** Count one Major issue for each missing Outcome rubric that checks for an explicit requirement in the prompt. This dimension only concerns "Outcome" criteria. Outcome rubrics must collectively check for: • Whether the user's request was fulfilled or not (e.g., was the email sent? was the investigation completed?) • Whether the final user-facing response is accurate and aligned with findings (e.g., correct names, dates, dollar amounts) • ~~Whether any required actions succeeded (e.g., tool returned success)~~ **(03/20)**

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

		Flag Missing Criteria — Outcome when: • The rubric set contains no criterion checking whether the user's core request was fulfilled • The rubric evaluates only process/trajectory behavior with no user-facing outcome checks • An explicit prompt requirement that affects the final answer or action is not represented by any Outcome criterion • A write action (send email, create Linear issue, update CRM, schedule calendar event, delete contact – basically any action that changes the state of the universe) has no Outcome rubric verifying it succeeded with correct details **(03/20)** • ~~An investigation-heavy task has no grounding-check rubric verifying the agent actually retrieved data before reporting it~~ **(03/20)**
	Criteria Not Self-Contained	**Definition:** Criterion cannot be evaluated against the model response without access to the universe, personas, tool schemas, and/or external facts/information. Every rubric must be self-contained. The LLM judge sees the trajectory and final response, all the rubric criteria (one rubric item can be used as context for another) and the original prompt — not the universe, not the personas, not the tool schema or any other external information. All expected values must be embedded directly in the criterion text. This applies to process/reasoning rubrics too. They can't be dependent on the tool outputs. The only exception is when more than one tool is valid for answering a prompt, in which case the rubrics should test for the intent of the model and can depend on tool outputs instead of being precise and self-contained. **Bad rubric examples (not self-contained):** • ~~[TS, Process] "The model must call the search tool with the year mentioned in the prompt." → Can't evaluate without knowing what year the prompt mentions, and the prompt is not in the reviewer's context.~~ • ~~[Outcome] "The model response must correctly summarize the output of the CRM tool call." → Can't evaluate without seeing the CRM tool output, and as an Outcome rubric, the tool output is not in the reviewer's context.~~ • [Outcome] "Email sent to the CEO." → Judge doesn't know who the CEO is. Must say: "Email sent to elena.rostova@moveops.com (CEO)." • [Outcome] "Budget is correct." → Must say: "Total cost is approximately $15,000 or less."

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

		• [Outcome] "The agent contacted the right person." → Must say: "Email sent to chloe.vance@moveops.com." (if "the right person" was not mentioned in the prompt)
		
		Good rubric examples (self-contained):
		
		• [Outcome] "An email was sent (via `send_email`) to chloe.vance@moveops.com." → Expected value fully embedded **(03/20)**
		  ○ **Note:** Write actions (any action that changes the state of the universe) are always Outcome only, not QC.
		• [Outcome] "The final response lists Noah Fitzgerald (GreenStack Solutions) as relocating to Seattle." → All expected values (name, company, city) are in the criterion.
		• [QC] "The model must pass a query related to 'Fatimah' or 'Al-Rashidi' (or similar) when calling `crm_search_contacts`." → Even though the rubric says "or similar," the intent is clear and evaluable **(03/20)**
		  ○ **Note:** Valid QC only if 'Fatimah' was not given in the prompt and the agent had to infer who to search for.

	Criteria Not Atomic	**Definition:** Criterion groups two or more constraints that are unrelated, which results in a rubric item with no clear focus on what aspect of the response it's trying to evaluate. Each rubric should evaluate one thing only — no bundling of multiple independent behaviors. Ask yourself: if this criterion fails, is there exactly one clear reason why? Acceptable bundling (NOT a violation): **QC rubrics may bundle multiple parameters of the same tool call:** "The model must pass 'fatimah.al-rashidi@moveops.com' as sender and "2026-02-14" as the before parameter when calling search_emails." Acceptable because parameters describe one action (searching for specific emails). The CB may choose either bundle or separate approach; both are acceptable. If parameters are missing from bundled QC criteria, count it as 1 missing criterion because it could be fixed with a single new/corrected criterion. QC rubrics may bundle multiple parameters of the same read/lookup tool call only — **never write actions (any action that changes the state of the universe is a write action)**. Write action parameters (recipients, content, dates) belong in Outcome rubrics 1.1/1.2, not QC **(03/20)**

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

		Outcome rubrics may bundle tightly coupled facts from the same source: "The final response lists Noah Fitzgerald (GreenStack Solutions) relocating to Seattle." Name + company + city are all attributes of the same relocation record and would be right or wrong together. General rule (03/04): Assessment of multiple components within one tool output can be grouped. Components from separate tool outputs should NOT be grouped. **Not atomic (violation):** "Email sent to Chloe and Linear issue created for Fatimah." → Independent actions — must be two separate rubrics. "The agent searched for Fatimah's relocations and sent a summary email to Elena." → Investigation step + write action = two independent claims. "The email mentions the storm, provides the new city, and includes flight details." → Three independent content requirements — should be three rubrics (unless all come from the same tool output).
	Incorrect Criteria	**Definition:** A rubric criterion checks for behavior that does not align with the prompt requirements, the Oracle Events, or valid logic. Count one Major issue for each criterion that: • Contradicts the prompt or Oracle Events (e.g., checks for the wrong recipient, wrong date, wrong entity) • Contains a factual error that's not grounded in the tool outputs or universe data • Requires behavior that would lead to an incorrect or misleading outcome • Is not an explicit requirement in the prompt and implementing it does not make the response better in any way **Important clarifications:** • Incorrect ≠ overly strict. A strict criterion is acceptable if it reflects required behavior. • Incorrect ≠ suboptimal. If a criterion is merely inefficient or awkward, consider Minor or Moderate issues instead. • Before classifying any issue as "Incorrect Criteria," check if a different, more specific error category would apply. For example, if a criterion is overly specific, count it as "Overly Specific Criteria" (Minor), not Incorrect.

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

		• Minor Wording Errors are NOT counted as Incorrect Criteria. Minor wording errors that might make the criterion potentially misleading can be flagged as non-fails. See [Purely Non-Failing Issues > Rubric Wording Errors]. It should still be possible to read the criterion with the proper intent; otherwise, this does not apply. **Auditor source of truth:** Use the Oracle Events / universe data as the source of truth. Universe data should be treated as static for evaluation purposes. **Incorrect vs Overly Broad:** If a criterion accepts all valid responses, but happens to also accept some invalid responses, flag that as [Overly Broad Criteria] **(03/26)**
Moderate Issues	**Missing Criteria — Process/Reasoning** **(Updated 03/20** - Significant updates to TS and QC scope, read everything again)	**Definition:** Count each missing process/reasoning rubric that represents a critical step in the Oracle Events as one issue. Critical steps are those without which you can't imagine a successful trajectory. This dimension only concerns: • Tool Selection (TS) • Query Construction (QC) **Default tools (always skip TS) — all read/lookup tools in Email, Slack, and Contacts** **Tool Selection guidance 📝:** • Write one TS rubric per non-default tool the agent might skip — one per tool regardless of how many times it's called • 📝 **Default tools** should **never** have TS rubrics — the agent always uses these, so TS provides no signal • Non-default tools (e.g., `crm_*`, `linear_*`, `calendar_*`, `airtable_*`, `quickbooks_*`) warrant TS if the agent might skip them entirely ○ **04/11 NOTE**: Non-default tools for **v2.2** universe: `mortgage_los_*, stripe_*, crm_*, filesystem_*, quickbooks_*`) • 📝 **Write actions** (`send_email`, `linear_create_issue`, `airtable_update_records`, `crm_update_deal_stage`,

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

`contacts_delete_contact`, etc.) **never** get TS or QC rubrics — they always go into Outcome. **A write action is any action that changes the state of the universe** (add, modify, delete as opposed to just reading)

Query Construction guidance 📝:

- 📝 Write QC only when **both** conditions hold: (1) the agent had to figure out the query itself — the value wasn't handed to it in the prompt, AND (2) a bad query would cause it to miss critical data entirely. If either condition fails, skip QC.
- 📝 QC is **never** written for write actions — write action parameters (recipients, dates, content, amounts) belong in Outcome rubrics (1.1/1.2), not QC

📝 **Process Rubric Decision Matrix:**

Situation	TS?	QC?	Example
Non-default tool + query value given in prompt	✅ Yes	❌ No	CRM lookup for "Noah Fitzgerald" — agent might skip CRM entirely, but name is given in prompt
Non-default tool + agent must infer or construct query AND bad query could miss critical data.	✅ Yes	✅ Yes	QuickBooks search for billing dispute — agent might skip AND might search "Heartland crating" instead of "Heartland billing"

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

Default tool + agent must infer or construct query AND bad query could miss critical data.	❌ No	✅ Yes	Slack search for buried financial discussion — agent always searches Slack, but narrow query misses the thread
Default tool + query value given in prompt or obvious	❌ No	❌ No	Email search for Elena Rostova — agent always searches email, and "Elena" or "Rostova" both return it

📝 Common critical steps that should typically have rubric coverage:

- **Tool Selection:** Each distinct non-default read/lookup tool in the Oracle Events should have one TS rubric. 📝 If a key non-default tool (e.g., `crm_search_contacts`, `linear_list_issues`, `quickbooks_search_invoices`) has no TS rubric, flag it. One TS criterion per tool even if called multiple times.
- **Query Construction:** Critical read/lookup tool calls where the agent had to construct a non-obvious query AND a bad query would miss critical data should have QC rubrics. Skip QC when the query value is given directly in the prompt or any reasonable keyword finds the data.

Flag Missing Criteria — Process when:

- A critical Oracle Event read/lookup step has no explicit rubric criterion
- The step is not reasonably encompassed by another existing criterion

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

		Do NOT flag missing process criteria for: • Write actions (any action that changes the state of the universe) — these belong in Outcome, not TS or QC • Default tools — TS is never required for these • Exploratory searches the agent might or might not make • Steps clearly and robustly covered by another criterion • Non-critical steps Notes: • Do not double-count missing steps if one criterion already reasonably covers them • Missing process criteria are Moderate, not Major, unless they completely prevent evaluation • Process criteria evaluate what actions occurred, not whether the final answer was correct
	Overlapping or Redundant Criteria **(Updated 03/20** - updated examples, read everything again)	**Definition:** Criterion is either completely redundant because other criteria completely encompass it, or multiple criteria check for the same thing partly. Count each completely redundant criterion as one Moderate issue, or count multiple overlapping criteria as one Moderate issue. **Redundant scenario:** • Criterion 1: "The model must use `crm_search_contacts` to find Fatimah's client records." • Criterion 2: "The model must use `crm_search_contacts` to find Fatimah's relocations." • Both are TS rubrics for the same tool — only one is needed regardless of the intent context. Per the one-TS-per-tool rule, write: "The model must use `crm_search_contacts`." **Example 2:** • Criterion 1: "An email was sent (via send_email) to chloe.vance@moveops.com." [Outcome 1.1] • Criterion 2: "The agent successfully contacted chloe.vance@moveops.com." [Outcome]

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

		• Both check the same thing — email sent to Chloe. One is redundant. **TS, QC, and Outcome do not overlap by design:** They operate on completely different dimensions — TS on read/lookup tool selection, QC on read/lookup parameters, Outcome on write action results and reported facts. A task can have all three without any overlap. **Notes:** • Key test: Would removing one criterion change scoring outcomes? If not, redundancy exists. • One TS rubric per non-default tool maximum — even if the tool appears in multiple OEs. • Never flag TS + QC + Outcome as overlapping when they each check a genuinely distinct dimension.

	Incorrectly Labeled Category
	
	(Updated 03/20 - significant changes, read everything again)
		Definition: Criterion is incorrectly categorized by selecting the wrong rubric category.
		
		Category List:
		
		• **Tool Selection:** Did the agent select the correct tool? **Read/lookup tools only — never write actions.**
		• **Query Construction:** Did the agent pass the correct parameters? **Read/lookup tools only — never write actions.**
		• **Outcome:** What was the result? What does the user see? All write actions go here. Outcome subtypes are:
		  ○ **[1.1]** Write-action results — did the right action happen with the right details? (verified from trajectory)
		  ○ **[1.2]** Action content — does the content match what was needed? (verified from trajectory parameters)
		  ○ **[2.1]** Key facts/findings — only when the final response is the deliverable (pure investigation tasks)
		
		Default tools (always skip TS) — all read/lookup tools in Email, Slack, and Contacts
		
		Common mislabeling errors:
		
		• A criterion checking specific parameter values labeled as Tool Selection → should be QC
		• A criterion checking whether an email succeeded labeled as QC → should be Outcome
		• A TS or QC rubric written for a write action (any action that changes the state of the universe) → should be Outcome. Write actions never appear in TS or QC.
		• A TS rubric written for a default tool → should be removed. Default tools never get TS rubrics.

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

		• A criterion that only checks which tool was used labeled as QC → should be TS (e.g., "The model must use `crm_search_contacts` to find Fatimah's relocations" is TS, not QC) **Clarification on TS vs QC boundary:** TS criteria check which tool was used, and may include intent context. Both of these are valid TS: • "The model must use the `crm_search_contacts` tool." • "The model must use a CRM search tool (e.g., `crm_search_contacts` or `crm_search_companies`)." QC criteria check specific parameter values passed to a read/lookup tool call, **but only when the agent had to figure out the query itself AND a bad query would miss critical data.** If the value was given in the prompt or any reasonable keyword finds the data, skip QC. ✅ Valid QC — agent had to infer: • "The model must pass a query related to 'Fatimah' or 'Al-Rashidi' (or similar) when calling `crm_search_contacts`." → Valid only if 'Fatimah' was NOT given in the prompt and the agent had to figure out who to search for from context. • "The model must pass 'completed' as the state filter when calling `linear_list_issues`." → Valid only if the agent had to determine the right status filter from context, and using the wrong filter would miss the relevant issues. ❌ Not valid QC — query given in prompt: • "The model must pass 'Noah Fitzgerald' when calling `crm_search_contacts`." → If 'Noah Fitzgerald' is named in the prompt, skip QC — agent can't meaningfully get this wrong. **The dividing line:** read/lookup tool + agent had to infer the query + bad query = missing data → QC. Query handed to agent in prompt → skip QC regardless of tool type. Write action + any parameter → Outcome.
	Unwanted Rubrics **ADDED 03/20**	**Definition:** • Tool selection or Query construction rubrics written for “write” actions are unwanted

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

		• Tool selection rubrics written for default tools are unwanted as these are always expected to be used by the agent. Adding these to the rubric doesn't make it worse, but they don't really add much value either, so CBs are instructed to skip these.
	Overly Broad Criteria **ADDED 03/26**	**Definition:** Criterion is too broad or overly permissive such that it would accept all valid responses, but also some invalid responses too **Example:** • Prompt: "Give me the contact details of exactly one point of contact for xyz client" • Correct answer: any of John, Jack or Jill • Overly Broad Criterion: "The model must mention one of John, Jack, Jill or Joe as the relevant point of contact for xyz client." • This criterion would accept three valid responses but also one invalid response (Joe) **Exception:** If the invalid paths that are accepted by the overly permissive criterion are unlikely to occur, do not flag the criterion as overly broad. See 69c2c02615c3291cfaa0d754 in the Examples tab for an explanation on this scenario.
Minor Issues	**Overly Specific Criteria** **(Updated 03/20** - updated examples and scenarios, read everything again)	**Definition:** Criterion is too specific such that it only accepts a subset of valid responses and would falsely punish the model for a completely valid actions. The criteria should not be overly specific in how the model can fulfill the goal. If there are multiple valid tools or routes the model can take, the criteria should not falsely punish the model for a valid decision. **Two main scenarios where this applies:** **1. Tool alternatives:** When more than one tool is valid for answering a prompt, the rubrics should test for the intent of the model. Examples: • Bad: "The model must use the `crm_search_contacts` tool." *(when `crm_search_companies` could also retrieve the same info)*

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

- Good: "The model must use a CRM lookup tool (e.g., `crm_search_contacts` or `crm_search_companies`)."

2. Freetext parameters: A QC criterion specifies an exact parameter value without allowing for alternatives when the tool accepts freetext.

- Overly specific: "The model must pass 'Fatimah relocations recent' to the query parameter when calling crm_search_contacts."
- Correct: "The model must pass a query related to 'Fatimah' or 'Al-Rashidi' (or similar) when calling crm_search_contacts."

3. Outcome content requirements: A 1.2 or 2.1 criterion specifies exact wording for agent-generated content when multiple valid expressions would satisfy the intent.

- ❌ Overly specific: "The email body must include the phrase 'we apologize for the inconvenience caused by the storm.'"
- ✅ Correct: "The email body includes an acknowledgment of the storm disruption (or similar)."
- ❌ Overly specific: "The final response must state 'Denver is the cheapest option at $11,200.'"
- ✅ Correct: "The final response identifies Denver as the lowest-cost option with an approximate cost figure (or similar)."

Exception — Strict values are never overly specific:

- Email addresses, IDs, dates, and exact strings pulled directly from universe data should always be exact. These are not overly specific — there is only one correct value.
- ✅ "The email body includes the client name 'GreenStack Solutions.'" — exact string from data, correct as strict.
- ✅ "An email was sent (via send_email) to chloe.vance@moveops.com." — one correct recipient, correct as strict.

Rules of thumb:

- Structured fields with one correct value (email addresses, IDs, dates) should be exact — NOT overly specific
- Freetext fields (search queries, email subjects, issue titles) must include "(or similar)" or list alternatives — exact values here ARE overly specific
- The test: could a competent agent pass a different but equally valid value? If yes, the criterion should allow it

QC criteria can be vague or nondescript if multiple queries could be considered valid to fulfill an intent. Example below:

[QC]: The model must pass a query related to 'weather API' or 'API outage' (or similar) when calling

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

		`conversations_search_messages`." → Any number of query strings could work — the focus is on the intent, not the exact string. Acceptable to leave vague when no specific query is required. **Note: Write action (any action that changes the state of the universe) parameters always go in Outcome — never TS or QC.**
Purely Non-Failing Issues	**Rubric Wording Errors**	**Rubric Wording Errors** **Definition:** Minor wording errors or typos in the rubric criteria that slightly alter the meaning of what's being checked, but in practice would not cause an LLM judge to evaluate incorrectly. **Example:** • Prompt asks about "Q3 travel spending" • Rubric says: "The final response must include Q3 travel expenses" • "Spending" vs "expenses" is a minor wording difference that wouldn't affect evaluation **Example:** • Prompt asks to check what "Friday special" items were available • Tool output lists "Sausage Roll" as a Friday special item • Rubric says: "The model must correctly identify sausage roll as one of today's specials" • "Friday special" vs "today's specials" is a minor wording error — semantically different, but in most cases wouldn't affect trajectory evaluation **How to assess:** Would an LLM judge, reading the trajectory, prompt and the rubric criteria, flag the response as incorrect because of the wording difference? If not, it's a non-failing wording error. If yes, escalate to Minor (Overly Specific) or Major (Incorrect) depending on impact. It should still be possible to read the criterion with the proper intent; otherwise, this exception does not apply and the issue should be escalated.

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z

		See [Major Issues > Incorrect Criteria] for more context.

List of “default tools” (03/20)

```
None
Slack: channels_list, conversations_history, conversations_replies, conversations_search_messages

Email: search_emails, list_emails, get_email_by_id, get_email_by_index, get_thread, get_email_threads, show_data,
list_jmap_emails, get_jmap_email, list_mailboxes, get_mailbox

Contacts: contacts_search_contacts, contacts_get_contacts, contacts_get_contact, contacts_show_data,
contacts_get_current_user_details
```

CONFIDENTIAL | latam.coder1946@remotasks.com+outlier | 2026-04-16T20:16:20.974Z